



Function Approximation

Blink Sakulkueakulsuk

Four main Bellman Equations

Bellman Expectation Equation

$$v_{\pi}(s) = \mathbb{E}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s, A_t \sim \pi(s)]$$

$$q_{\pi}(s, a) = \mathbb{E}[R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1}) | S_t = s, A_t = a]$$

Bellman Optimality Equation

$$v^*(s) = \max_a \mathbb{E}[R_{t+1} + \gamma v^*(S_{t+1}) | S_t = s, A_t = a]$$

$$q^*(s, a) = \mathbb{E} \left[R_{t+1} + \gamma \max_{a'} q^*(S_{t+1}, a') \mid S_t = s, A_t = a \right]$$

Prediction and Control

Prediction

- A problem when we perform **policy evaluation**
- Estimating v_π or q_π

Control

- A problem when we perform **policy optimization**
- Estimating v^* or q^*

Policy Evaluation

Given a policy, we want to estimate this

$$v_{\pi}(s) = \mathbb{E}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) | s, \pi]$$

We initialize all value to zero, and we iterate the equation above as an update

$$v_{k+1}(s) = \mathbb{E}[R_{t+1} + \gamma v_k(S_{t+1}) | s, \pi] \quad , \quad \forall s$$

Note that

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_r \sum_{s'} p(r, s' | s, a) (r + \gamma v_{\pi}(s'))$$

If $v_{k+1}(s) = v_k(s), \forall s$, we solve v_{π} .

Policy Improvement

The example shows that we can improve a policy when we learn the value.

- In the example, we do not improve any policy

$$\begin{aligned}\forall s: \pi_{new}(s) &= \operatorname{argmax}_a q_{\pi}(s, a) \\ &= \operatorname{argmax}_a \mathbb{E}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s, A_t = a]\end{aligned}$$

We use the new policy to evaluate, and repeat.

Monte-Carlo Policy Evaluation

Given the sequential decision problem, MC learns v_π from episodes under policy π

$$S_1, A_1, R_2, \dots, S_k \sim \pi$$

The return is calculated given an ending time T

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-t-1} R_T$$

The value function is then calculated from the expected return

$$v_\pi(s) = \mathbb{E}[G_t | S_t = s, \pi]$$

We can use **sample average return** instead of **expected return**

- This is called **Monte Carlo policy evaluation**

Temporal Difference Learning

$$v_{k+1}(s) = \mathbb{E}[R_{t+1} + \gamma v_k(S_{t+1}) | S_t = s, A_t \sim \pi(S_t)]$$

Instead of calculating the expected return, we can sample the value.

$$v_{t+1}(S_t) = R_{t+1} + \gamma v_t(S_{t+1})$$

But instead of updating everything on the noisy value, we can update the value a little bit instead.

$$v_{t+1}(S_t) = v_t(S_t) + \alpha_t(R_{t+1} + \gamma v_t(S_{t+1}) - v_t(S_t))$$

MC Prediction vs TD Prediction

Prediction problem: learn v_π online from experience under policy π

Monte Carlo:

- Update value $v_n(S_t)$ with respect to sample return G_t

$$v_n(S_t) = v_n(S_t) + \alpha(G_t - v_n(S_t))$$

Temporal-difference learning:

- Update value $v_t(S_t)$ with respect to estimated return $R_{t+1} + \gamma v(S_{t+1})$

$$v_{t+1}(S_t) = v_t(S_t) + \alpha_t(R_{t+1} + \gamma v_t(S_{t+1}) - v_t(S_t))$$

Bootstrapping and Sampling

Bootstrapping

- Use the estimate of the next state to update
- DP and TD use
- MC does not

Sampling

- Update samples an expectation
- MC and TD sample
- DP does not

Temporal Difference Learning

We can also learn action values by updating value $q_t(S_t, A_t)$ with respect to estimated return $R_{t+1} + \gamma v(S_{t+1}, A_{t+1})$

$$q_{t+1}(S_t, A_t) = q_t(S_t, A_t) + \alpha_t(R_{t+1} + \gamma q_t(S_{t+1}, A_{t+1}) - q_t(S_t, A_t))$$

This algorithm is called SARSA, because it uses $S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1}$

Multi-Step Returns

In general, we can formalize n-step return as

$$G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} v(S_{t+n})$$

We can define multi-step temporal-difference learning as

$$v(S_t) \leftarrow v(S_t) + \alpha(G_t^{(n)} - v(S_t))$$

Mixing Multi-Step Returns


Consider a multi-step return equation,

$$G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} v(S_{t+n})$$

Multi-step returns bootstrap on one state, $v(S_{t+n})$. However, we can modify the equation so you can bootstrap a little bit on many states

$$G_t^\lambda = R_{t+1} + \gamma((1 - \lambda)v(S_{t+1}) + \lambda G_{t+1}^\lambda)$$

Monte Carlo Control



Perform policy evaluation


- Using Monte Carlo policy evaluation

$$q \approx q_\pi$$

Perform policy improvement

- Using **ϵ -greedy** policy improvement

SARSA



Perform policy evaluation

- Using SARSA

$$q \approx q_\pi$$

Perform policy improvement

- Using ϵ -greedy policy improvement

On-Policy and Off-policy Learning

On-policy learning

- Learn from actually performing
- Learn a policy π from samples from π

Off-policy learning

- Learn from other sources
- Learn a policy π from samples from b

Double Q-learning

Since

$$\max_a q_t(S_{t+1}, a) = q_t(S_{t+1}, \operatorname{argmax}_a q_t(S_{t+1}, a))$$

Q-learning will use the same values to select action, and evaluate the value

$$R_{t+1} + \gamma \max_{a'} q_t(S_{t+1}, a') = R_{t+1} + \gamma q_t(S_{t+1}, \operatorname{argmax}_a q_t(S_{t+1}, a))$$

And this will result in the issue in the last slide. To solve this problem, we can introduce another policy to split action selection from evaluation.

Double Q-learning

Double Q-learning will

- store two q functions: q and q'
- update one q function for each iteration.

$$\begin{aligned}
 &R_{t+1} + \gamma q'_t(S_{t+1}, \operatorname{argmax}_a q_t(S_{t+1}, a)) \\
 &R_{t+1} + \gamma q_t(S_{t+1}, \operatorname{argmax}_a q'_t(S_{t+1}, a))
 \end{aligned}
 \left. \vphantom{\begin{aligned} R_{t+1} + \gamma q'_t(S_{t+1}, \operatorname{argmax}_a q_t(S_{t+1}, a)) \\ R_{t+1} + \gamma q_t(S_{t+1}, \operatorname{argmax}_a q'_t(S_{t+1}, a)) \end{aligned}} \right\} \begin{array}{l} \text{เลือกอย่างใดอย่างหนึ่ง} \\ \text{ในการอัปเดตแต่ละรอบ} \end{array}$$

Off-policy Q-Learning Control

To control, we want to optimize both policies.

For the target policy π , we greedify the action on $q(s, a)$

$$\pi(S_{t+1}) = \operatorname{argmax}_{a'} q(S_{t+1}, a')$$

For the behavior policy b , we can use ϵ -greedy on $q(s, a)$

all data of 6 weeks before is clear

Thus, Q-Learning target is simplified as:

but in big problem we cannot

$$R_{t+1} + \gamma \sum_a \pi(a|S_{t+1}) q(S_{t+1}, a) = R_{t+1} + \gamma \max_a q(S_{t+1}, a)$$

Function Approximation

If we look closely, there are a lot of functions in many parts of reinforcement learning.

- policy a function that outputs an action given a state
- value function a function that outputs a value of a state
- agent state update a function that outputs a next state given some inputs
- model a function that outputs dynamic of the environment given some inputs

Function Approximation

All the functions are what we want to learn, but what if the problem space is too large?

- There are too many states, or too many actions

OR continuous 6 weeks before is useless. focus this week!!!

We can approximate these functions.

- This is often called **deep reinforcement learning**, when we use **deep neural network** to represent these functions.

Function Approximation and Large Problem Space

In recent years, deep reinforcement learning is used to solve large problems.

before dynamic programming can solve medium problems.

- Go: 10^{170} states
- Drone: continuous state space
- Robots: real world space

Value Function Approximation

๖ ๓๐๓๓๓๓๓๓๓๓๓๓๓๓๓๓๓๓ array or sth.

So far, we use **lookup tables** to store all values.

- Storing a value $v(S)$ for each state S
- Storing an state-action value $q(S, A)$ for each pair of S, A

In a large problem, this method will run into problems.

- Too many states to store in memory
- Too slow to learn values of states individually
- States may not be fully observable

Value Function Approximation

To solve large MDPs, we can estimate the value function using function approximation instead

W - matrix, vector
 model ML model

function Approximation $v_{\mathbf{w}}(s) \approx v_{\pi}(s)$ look up table

$q_{\mathbf{w}}(s, a) \approx q_{\pi}(s, a)$

ค่าฟังก์ชันของสถานะ: ค่าฟังก์ชันของ
 คู่สถานะ-การกระทำ และค่าฟังก์ชันของสถานะ-การกระทำ
 supervised learning

- We can generalize new states (unexplored) from old states (explored)
- We can use MC or TD to learn parameter \mathbf{w} .

Agent State Approximation

If we encounter a non-fully observable problem, we can approximate agent state instead.

agent state transition

$$\mathbf{s}_t = u_{\omega}(\mathbf{s}_{t-1}, A_{t-1}, O_t) ; \text{update } \mathbf{s}_t$$

input vector can represent State

parameter

vector

$\omega \in \mathbb{R}^n$ denotes learnable parameters

\mathbf{s}_t denotes the agent state

- Think of this as a vector of features in a state

Function Approximation Classes

class 1

Tabular: A table containing states

State aggregation: Partitioning environment states into discrete states

↳ transition เป็น ร่อง รวมห้อง and
represent ในหัวข้อเป็น ตัวอักษร

class 2

Linear function approximation

- Fixed agent state update → fixed dimension
- Fixed feature map $\mathbf{x} : \mathcal{S} \Rightarrow \mathbb{R}^n \rightarrow$ fixed feature
- Linear function as value function approximation $v_{\mathbf{w}}(s) = \mathbf{w}^T \mathbf{x}(s)$

class 3

Differentiable function approximation

- $v_{\mathbf{w}}(s)$ is a differentiable function. It can be non-linear.
- We can utilize Deep Learning to learn features, such as CNN.

Function Approximation Classes

Generally, any function approximator works. But reinforcement Learning has some specific properties too keep in mind.

✍ มีประสบการณ์ที่ต่างกัน และ independent

- Experience is not independent and identically distributed (i.i.d.). Current and future actions are correlated.
- Policy affects the observations and rewards
- Learning target can be non-stationary

✍ target เปลี่ยนไปเมื่อ policy

Policy change, bootstrapping, dynamic of the problem

Function Approximation

How to choose what to use? * ใช้ไม่ได้ทุกอย่างหมด

Tabular

Good theory, not scalable

Linear function approximation

Good theory, need good feature construction

Non-linear function approximation

Scalable, but not many theories to back up (yet)

Less reliable on picking good features

Gradient Descent Revisited

Given $J(\mathbf{w})$, a differentiable function of parameter vector \mathbf{w}

↓ is gradient

The gradient of $J(\mathbf{w})$ is

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \begin{pmatrix} \frac{\partial J(\mathbf{w})}{\partial \mathbf{w}_1} \\ \vdots \\ \frac{\partial J(\mathbf{w})}{\partial \mathbf{w}_n} \end{pmatrix}$$

Gradient Descent Revisited

$$\nabla_{\mathbf{w}} J(\mathbf{w}) = \begin{pmatrix} \frac{\partial J(\mathbf{w})}{\partial \mathbf{w}_1} \\ \vdots \\ \frac{\partial J(\mathbf{w})}{\partial \mathbf{w}_n} \end{pmatrix}$$

We want to find the minimum value of $J(\mathbf{w})$. To do that, we can update \mathbf{w} such that it moves in the direction of negative gradient

$$\Delta \mathbf{w} = -\frac{1}{2} \alpha \nabla_{\mathbf{w}} J(\mathbf{w})$$

α is learning rate or step-size parameter.

Approximation using Stochastic Gradient Descent

* point is direction not exact value *

Given an approximator with parameters \mathbf{w} , we want to minimize the value difference between this approximator $v_{\mathbf{w}}(S)$ and true value $v_{\pi}(S)$

$$J(\mathbf{w}) = \mathbb{E}_{S \sim d} [(v_{\pi}(S) - v_{\mathbf{w}}(S))^2]$$

\mathbb{E} : expectation (คาดหวัง sth)
 d : distribution of states
 $(v_{\pi}(S) - v_{\mathbf{w}}(S))^2$: error² (constant)
 $v_{\mathbf{w}}(S)$: learn w ที่ทำให้ $v_{\mathbf{w}}(S)$ ใกล้เคียง

d denotes a distribution of states

So the gradient descent is:

$$\Delta \mathbf{w} = -\frac{1}{2} \alpha \nabla_{\mathbf{w}} J(\mathbf{w}) = \alpha \mathbb{E}_d [(v_{\pi}(S) - v_{\mathbf{w}}(S)) \nabla_{\mathbf{w}} v_{\mathbf{w}}(S)]$$

α : learning rate (อัตรา)
 $\frac{1}{2}$: cancel diff
 \mathbb{E}_d : expectation (คาดหวัง)
 $\nabla_{\mathbf{w}} v_{\mathbf{w}}(S)$: diff (ดิฟ)
 $v_{\mathbf{w}}(S)$: discrete state (ไม่ต่อเนื่อง ไม่ได้นับ) (ดู next slide)
 ∇ : sample 1 ep. meaning (ตัวอย่าง)

In sampling algorithm (MC, TD), **Stochastic** gradient descent can be used to sample the gradient.

$$\Delta \mathbf{w} = \alpha (v_{\pi}(S_t) - v_{\mathbf{w}}(S_t)) \nabla_{\mathbf{w}} v_{\mathbf{w}}(S_t)$$

Feature Vectors

If we want to approximate some functions, generally, it relies on states. To make it computable, we represent a state as a **feature vector**

$$\begin{array}{c} \text{fixed mapping} \\ \text{state to feature} \\ \uparrow \\ \mathbf{x}(s) = \begin{pmatrix} x_1(s) \\ \vdots \\ x_n(s) \end{pmatrix} \end{array}$$

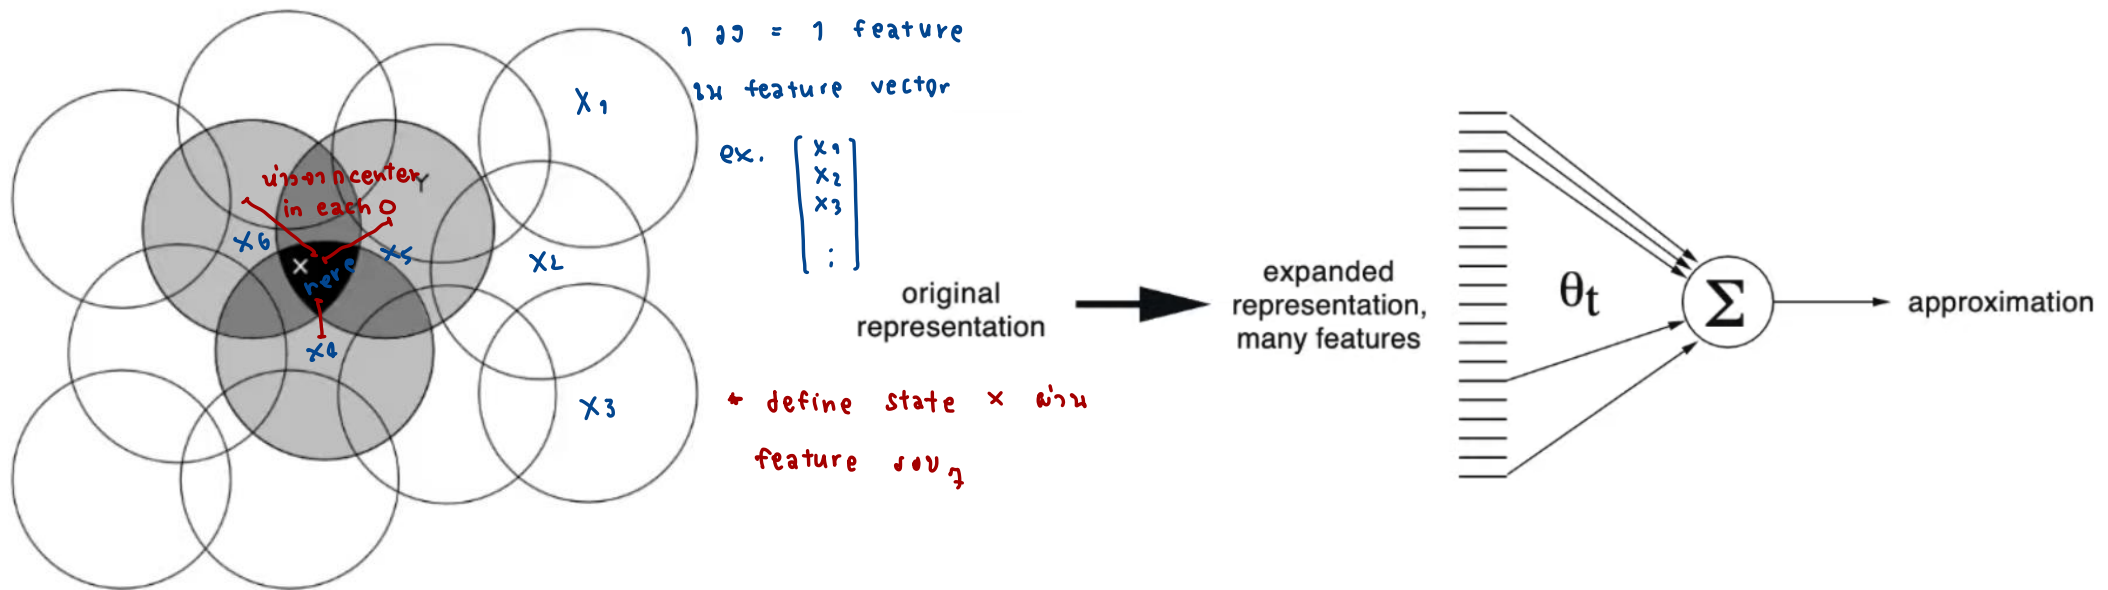
$\mathbf{x}: S \rightarrow \mathbb{R}^n$ is a fixed mapping from state to features

We can represent a feature vector of state S_t as $\mathbf{x}_t = \mathbf{x}(S_t)$

- Distance of a robot from reference points
- Objects in an image

Coarse Coding วิธีสร้าง feature vector

Coarse coding is a method to construct a feature vector through state representation.



Linear Value Function Approximation

Let's consider a simple case, a linear function approximator. We can express the value function as

$$v_{\mathbf{w}}(S) = \underbrace{\mathbf{w}^T}_{\substack{\text{weight} \\ \text{feature}}} \underbrace{\mathbf{x}(S)}_D = \sum_{j=1}^n \mathbf{x}_j(S) \mathbf{w}_j$$

The loss of this value function is then defined as

$$J(\mathbf{w}) = \mathbb{E}_{S \sim d} \left[\left(v_{\pi}(S) - \mathbf{w}^T \mathbf{x}(s) \right)^2 \right]$$

Linear Value Function Approximation

$$J(\mathbf{w}) = \mathbb{E}_{S \sim d} \left[\left(v_{\pi}(S) - \mathbf{w}^T \mathbf{x}(S) \right)^2 \right]$$

$$v_{\mathbf{w}}(x) = w_1 x_1 + w_2 x_2 + \dots$$

$$\nabla v = \left[\frac{\partial (v_{\mathbf{w}}(x))}{\partial w_1} \right] = \begin{bmatrix} x_1 \\ \vdots \end{bmatrix} = \vec{x}$$

$$\frac{\partial v_{\mathbf{w}}(x)}{\partial w_n} = \vec{x}$$

Then we can use stochastic gradient descent to update the parameter

- Stochastic gradient descent is a gradient descent algorithm that calculates loss of 1 input at a time
- We can update the parameter as

$$\nabla_{\mathbf{w}} v_{\mathbf{w}}(S_t) = \mathbf{x}(S_t) = \mathbf{x}_t \quad \Rightarrow \quad \Delta \mathbf{w} = \alpha (v_{\pi}(S_t) - v_{\mathbf{w}}(S_t)) \mathbf{x}_t$$

Update = step-size \times prediction error \times feature vector

Learning Targets

$$\Delta \mathbf{w}_t = \alpha (v_\pi(S_t) - v_{\mathbf{w}}(S_t)) \nabla_{\mathbf{w}} v_{\mathbf{w}}(S_t)$$

We cannot update $v_{\mathbf{w}}(S_t)$ toward $v_\pi(S_t)$ since we do not know the true value. However, we can change the target according to algorithms we use.

For MC, we use the return G_t

$$\Delta \mathbf{w}_t = \alpha \left(\overset{\substack{\text{return (total reward) / 1 episode} \\ \uparrow}}{G_t} - \underset{\substack{\text{learnable param} \\ \downarrow}}{v_{\mathbf{w}}(S_t)}} \right) \nabla_{\mathbf{w}} v_{\mathbf{w}}(S_t)$$

For TD, we use TD target $R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1})$

$$\Delta \mathbf{w}_t = \alpha (R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1}) - v_{\mathbf{w}}(S_t)) \nabla_{\mathbf{w}} v_{\mathbf{w}}(S_t)$$

For TD(λ),

$$\begin{aligned} \Delta \mathbf{w}_t &= \alpha (G_t^\lambda - v_{\mathbf{w}}(S_t)) \nabla_{\mathbf{w}} v_{\mathbf{w}}(S_t) \\ G_t^\lambda &= R_{t+1} + \gamma((1 - \lambda)v_{\mathbf{w}}(S_{t+1}) + \lambda G_{t+1}^\lambda) \end{aligned}$$

Monte-Carlo with Value Function Approximation

The return of MC is an unbiased, noisy sample of $v_\pi(s)$

We can use (online) supervised learning to train the model

$$\{(S_0, G_0), \dots, (S_t, G_t)\}$$

(trend like dataset
many episode

So, the linear MC updates is

plugin

$$\Delta \mathbf{w} = \alpha (G_t - v_{\mathbf{w}}(S_t)) \nabla_{\mathbf{w}} v_{\mathbf{w}}(S_t) = \alpha (G_t - v_{\mathbf{w}}(S_t)) \mathbf{x}_t$$

MC with linear approximation converges to the global optimum

MC with non-linear value function approximation also converges (to local or global optimum)

TD Learning with Value Function Approximation

The target of TD is a biased sample of $v_\pi(S_t)$


We can still use supervised learning to train the model

$$\{(S_0, R_1 + \gamma v_{\mathbf{w}}(S_1)), \dots, (S_t, R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1}))\}$$

So, the **linear** TD updates is

$$\Delta \mathbf{w}_t = \alpha (R_{t+1} + \gamma v_{\mathbf{w}}(S_{t+1}) - v_{\mathbf{w}}(S_t)) \nabla_{\mathbf{w}} v_{\mathbf{w}}(S_t) = \alpha \delta_t \mathbf{x}_t$$

Control with Value Function Approximation



Perform policy evaluation

- Approximate policy evaluation

$$q_w \approx q_\pi$$

Perform policy improvement

- Using **ϵ -greedy** policy improvement

Action-Value Function Approximation

We can apply function approximation to action-value function $q_{\mathbf{w}}(s, a) \approx q_{\pi}(s, a)$

For example, the linear action-value function is

$$q_{\mathbf{w}}(s, a) = \mathbf{x}(s, a)^T \mathbf{w}$$

Stochastic gradient descent update is

$$\begin{aligned} \Delta \mathbf{w} &= \alpha (q_{\pi}(s, a) - q_{\mathbf{w}}(s, a)) \nabla_{\mathbf{w}} q_{\mathbf{w}}(s, a) \\ &= \alpha (q_{\pi}(s, a) - q_{\mathbf{w}}(s, a)) \mathbf{x}(s, a) \end{aligned}$$

Deep Reinforcement Learning

We can use deep neural networks as our function approximator

- Apply DNN to approximate value function or action-value function
- Using DNN with MC, TD, Q, or Double-Q

OMG

Online Neural Q-Learning

Neural network: $O_t \rightarrow \mathbf{q}_w$

Exploration policy: $\pi_t = \epsilon\text{-greedy}(\mathbf{q}_t)$, then $A_t \sim \pi_t$

Weight update:

$$\Delta \mathbf{w} = k \left(R_{t+1} + \gamma \max_a q_w(S_{t+1}, a) - q_w(S_t, A_t) \right) \nabla_{\mathbf{w}} q_w(S_t, A_t)$$

Optimizer: SGD, Adam, ...

train by

Deep Q-Network

Neural network: $O_t \rightarrow \mathbf{q}_w$

Exploration policy: $\pi_t = \epsilon\text{-greedy}(\mathbf{q}_t)$, then $A_t \sim \pi_t$

A replay buffer to store and sample past transitions $(S_i, A_i, R_{i+1}, S_{i+1})$

Target network parameters \mathbf{w}^- \downarrow
for feature

Weight update:

$$\Delta \mathbf{w} = \left(R_{t+1} + \gamma \max_a q_{\mathbf{w}^-}(S_{t+1}, a) - q_{\mathbf{w}}(S_t, A_t) \right) \nabla_{\mathbf{w}} q_{\mathbf{w}}(S_t, A_t)$$

Update $\mathbf{w}_t^- \leftarrow \mathbf{w}_t$ periodically (i.e., every 1000 steps)

Optimizer: SGD, Adam,...